

1. What is HTTP Caching & Why is it needed?

HTTP Caching is a mechanism where the browser (and intermediary servers) stores a local copy of server responses.

- **The Need:** It reduces **latency** (faster loads), decreases **bandwidth** consumption, and lowers the **server load**.
- **The Performance Impact:** It is the difference between a **200 OK** (fetching from network, ~200ms-2s) and a **200 OK (from cache)** or **304 Not Modified** (~10ms-50ms).

2. Setting and Accessing: The Roles

- **Who Sets it?** The **Server** defines the caching policy via HTTP Response Headers (e.g., Cache-Control).
- **Who Accesses it?** The **Browser** (automatically). JavaScript/Frontend code **cannot** directly read or write to the HTTP Cache (unlike the Cache API or LocalStorage). It is managed by the browser's network stack.

3. Storage & Memory Location

- **Where it lives:**
 1. **Memory Cache (RAM):** Very fast, volatile. Stores resources used during the current session. Cleared when the tab/browser closes.
 2. **Disk Cache (Hard Drive):** Slower but persistent. Stores resources intended for long-term use.
- **Capacity:** Browsers typically allocate a percentage of available disk space (often hundreds of MBs to GBs). It is much larger than LocalStorage (5MB).

4. Eviction: How is it cleared?

- **Automatic:** The browser uses an **LRU (Least Recently Used)** algorithm. When the cache is full, the oldest/least accessed files are deleted.
 - **Manual (Client):** User clears "Browsing Data."
 - **Programmatic (Server):** The server cannot "reach into" a browser and delete a file. It can only "invalidate" it by changing the file's URL (Versioning/Hashing) or sending a Clear-Site-Data: "cache" header.
-

5. Various Ways to Set HTTP Cache (The Core Interview Topic)

There are two main categories: **Strong Caching** and **Validation (Consultative) Caching**.

A. Strong Caching (No Network Request)

The browser doesn't even talk to the server. It just grabs the file from the disk.

1. **Cache-Control (Modern/Recommended):**
 - o max-age=31536000: Cache the file for 1 year (in seconds).
 - o immutable: Tells the browser the file will *never* change (used with hashed filenames like main.a7b2.js).
 - o no-store: Do not cache at all (used for sensitive data).
 - o no-cache: This is a "gotcha." It **does not** mean "don't cache." It means "Cache it, but check with the server before using it" (Validation).
2. **Expires (Legacy):**
 - o Sets a specific date: Expires: Wed, 21 Oct 2025 07:28:00 GMT.
 - o **Problem:** If the user's system clock is wrong, this fails. Cache-Control overrides this.

B. Validation Caching (Network Request + 304 Response)

The browser asks: "I have this file from yesterday, is it still good?" If yes, the server sends a **304 Not Modified** (no body, very fast).

1. **ETag (Entity Tag):**
 - o A unique fingerprint/hash of the file content.
 - o **Flow:** Server sends ETag: "v1.2". Next time, Browser sends If-None-Match: "v1.2".
2. **Last-Modified:**
 - o A timestamp.
 - o **Flow:** Server sends Last-Modified: [Date]. Next time, Browser sends If-Modified-Since: [Date].

6. Critical Concepts You Must Mention

Cache-Busting (Fingerprinting)

Since you might set a file to cache for 1 year (max-age=31536000), how do you update it when you fix a bug?

- **The Strategy:** You change the filename. Instead of style.css, you use style.v2.css or style.8h2f.css.
- **Why?** The browser sees a new URL and treats it as a completely new resource, bypassing the old cache.

Private vs. Public

- Cache-Control: public: Can be cached by the browser **and** CDN/Proxies (good for static assets).
- Cache-Control: private: Only the **user's browser** can cache it (good for user-specific HTML).

Vary Header

- Tells the cache that the response depends on a specific request header (usually Vary: Accept-Encoding to distinguish between Gzip and Brotli versions of a file).
-

7. How it helps Performance: The "Metric" View

- **TTFB (Time to First Byte):** Reduced significantly because cached hits don't wait for server processing.
 - **LCP (Largest Contentful Paint):** Images (like hero banners) that are cached render almost instantly on the second visit.
 - **Bandwidth Cost:** Drastically reduces data costs for users on mobile plans and reduces your cloud egress costs.
-

8. Most Asked Interview Q&A

Q: "What is the difference between no-cache and no-store?"

A: no-store is absolute—the browser must not keep any copy. no-cache allows the browser to keep a copy but forces it to validate with the server (via ETag) before showing it to the user.

Q: "Explain the 'Stale-While-Revalidate' (SWR) header."

A: This is a modern performance strategy. Cache-Control: max-age=60, stale-while-revalidate=3600.

The browser shows the cached (stale) version immediately if it's less than an hour old, but simultaneously triggers a background fetch to update the cache for the next visit.

Q: "If a response has both Cache-Control: max-age=10 and an ETag, what happens after 15 seconds?"

A: The max-age has expired. The browser will make a request to the server, sending the If-None-Match (ETag) header. If the file hasn't changed, the server returns 304, and the browser reuses the local file.

1. What exactly can be cached?

The HTTP Cache is designed for **Idempotent** requests (requests that don't change the state of the server).

- **Methods:** Primarily **GET** requests. While some browsers theoretically allow caching of HEAD or POST (if configured with explicit freshness headers), in 99% of production systems, **only GET requests are cached**.
 - **Content Types:**
 - **Static Assets:** JS files, CSS files, Images (PNG, WebP, SVG), Fonts (WOFF2).
 - **Documents:** HTML files.
 - **Data:** API responses (JSON/XML) — though this is usually managed with shorter TTLs (Time to Live).
 - **What NOT to cache:** POST, PUT, DELETE requests, and any URL containing highly sensitive PII (Personally Identifiable Information) unless using Cache-Control: private.
-

2. The Flow: Strong Caching (Cache-Control: max-age)

In this scenario, the browser is the "Decision Maker."

Scenario A: Cache Hit

- The User requests app.js.
- The Browser checks its internal Cache Map.
- **Check:** Is `current_time < (received_time + max_age)`?
- **Result: Yes.**
- **Action:** The Browser retrieves the file from **Disk/Memory**. No network request is ever sent.
- **DevTools:** You see Status: 200 OK (from disk cache).

Scenario B: Cache Miss (Expired or Missing)

- The User requests app.js.
 - The Browser checks the cache.
 - **Result:** File is missing **or** `max_age` has expired.
 - **Action:** The Browser sends a full HTTP GET request to the Server.
 - **Server Response:** The Server sends the file (200 OK) + a new Cache-Control: `max-age=...` header.
 - **Update:** Browser saves the new version to the disk.
-

3. The Flow: Validation Caching (ETag / If-None-Match)

In this scenario, the Server is the "Decision Maker."

Scenario A: Cache Hit (Revalidation)

- The User requests style.css.
- The Browser sees it has a version but it's "stale" (expired max-age or no-cache was used). It sees an ETag: "xyz123" in its records.
- **Action:** The Browser sends a request with a header: If-None-Match: "xyz123".
- **Server Check:** The Server checks the current file hash. It still matches "xyz123".
- **Result:** The Server sends a **304 Not Modified** (no body/payload).
- **Action:** The Browser updates the "freshness" of its local copy and displays it. This is extremely fast because the response body is empty.

Scenario B: Cache Miss

3. The User requests style.css with If-None-Match: "xyz123".
4. **Server Check:** The file has changed! The new hash is "abc789".
5. **Result:** The Server sends a **200 OK** + the **entire new file** + the new ETag: "abc789".
6. **Action:** The Browser replaces the old file in the cache with the new one.

4. Summary Table for the Interviewer

Feature	Strong Cache (max-age)	Validation Cache (ETag)
Network Request?	No (if fresh)	Yes (always, to check)
Server Load	Zero	Minimal (Header processing only)
Use Case	Hashed assets (main.8h2f.js)	Mutable files (index.html)
Response Code	200 OK (from cache)	304 Not Modified

5. Pro-Tip: The "Hashed Asset" Pattern

In modern Frontend System Design, we combine these.

- **HTML (index.html):** Use Cache-Control: no-cache + ETag. We always check the server to see if a new version of the app exists.
 - **Assets (bundle.js, logo.png):** Use **Content Hashing** (e.g., main.a1b2.js). Set Cache-Control: max-age=31536000, immutable.
 - **The Benefit:** If the JS changes, the HTML will point to a new filename. The old file stays in cache but is ignored; the new file is fetched and cached forever.
-

This is a great clarifying question. In an interview, knowing the **default behavior** shows you understand the "magic" that happens under the browser's hood.

1. Which types of responses get cached?

Technically, the browser looks at the **HTTP Method** and the **Status Code**.

- **HTTP Methods:** Only **GET** is cached by default. POST, PUT, PATCH, and DELETE are considered "unsafe" or "non-idempotent" (they change data on the server), so the browser will **never** cache them by default.
- **Status Codes:** The browser typically only caches "Successful" responses.
 - **Commonly Cached:** 200 (OK), 203 (Non-Authoritative Info), 204 (No Content), 300 (Multiple Choices), 301 (Moved Permanently), 410 (Gone).
 - **Rarely/Never Cached:** 404 (Not Found), 500 (Internal Server Error), 503 (Service Unavailable).

2. Heuristic Caching: What happens if no headers are specified?

If the server sends a GET response but **does not** include Cache-Control, Expires, ETag, or Last-Modified, the browser doesn't just give up. It performs **Heuristic Caching**.

The browser thinks: "*The server didn't tell me what to do, but this looks like a static file. I'll guess how long it's safe to keep.*"

The "10% Rule" (Standard Heuristic)

Most modern browsers (Chrome, Firefox) use a simple formula if they see a Last-Modified header but no Cache-Control:

$$\text{Cache Duration} = (\text{Date of Request} - \text{Last-Modified Date}) * 0.10$$

- **Example:** If you fetch a file today (Jan 6, 2026) and the server says it was last modified 100 days ago, the browser will cache it for **10 days** automatically.
- **The Danger:** If the server provides **no headers at all** (not even Last-Modified), the

browser behavior is "implementation-defined." Usually, it won't cache it for long, or it might treat it as a session-only cache.

3. The Flow: Cache Hit vs. Cache Miss

To make this crystal clear for your interview, here is the step-by-step logic for both **Strong** and **Validation** caching.

Scenario 1: Strong Caching (The "Speed" Path)

This uses Cache-Control: max-age=....

- **Cache HIT:**
 - Browser: "I need hero.jpg."
 - Browser checks cache: "I have it. It was fetched 1 minute ago. Its max-age is 1 hour."
 - **Result:** Browser takes it from disk. **Network activity: 0.**
- **Cache MISS:**
 - Browser: "I need hero.jpg."
 - Browser checks cache: "It's expired (stale)" or "I don't have it."
 - **Result:** Browser sends a full request to the Server. Server returns 200 OK + the file.

Scenario 2: Validation Caching (The "Integrity" Path)

This uses ETag or Last-Modified.

- **Cache HIT (Revalidation):**
 - Browser: "I need index.html. My version is stale, but it has ETag: 'v1'."
 - Browser sends: GET /index.html with header If-None-Match: 'v1'.
 - Server: Checks its file. It's still 'v1'.
 - **Result:** Server sends **304 Not Modified** (no body). Browser uses its local copy.
- **Cache MISS:**
 - Browser sends: GET /index.html with If-None-Match: 'v1'.
 - Server: "The file changed. The new version is 'v2'."
 - **Result:** Server sends **200 OK** + the entire new index.html file + ETag: 'v2'.

4. Summary of "No Header" Consequences

In an FE System Design interview, if you are asked about a site with no cache headers, your answer should be:

2. **Inconsistency:** Different browsers will guess different durations (Heuristic Caching).
3. **Performance Risk:** A browser might cache a file for too long, and the user won't see updates (Stale Data).
4. **Server Stress:** Or, the browser might never cache it, causing the server to be hit with requests for every single image and icon on every page load.